

Programming Skills coursework II b: benchmarking percolate v0.1.0 with different densities

Abstract

Keywords:

1. Introduction

This paper documents the results of a benchmark performed on **percolate v0.1.0**. **percolate** is a scientific program written in the Fortran programming language. It generates a random matrix with two kinds of cells: empty and filled. Empty cells build clusters with their neighbors and **percolate** finds clusters that begin at the first column and end at the last. If such a cluster exists, the matrix percolates.

One way to configure **percolate** is by setting the density of the empty cells. This benchmark looks at the behavior of **percolate** with different densities of filled cells and how the density influences the execution time. It consists of two phases. During the first phase of the benchmark, the execution time of **percolate** with different densities is described and analyzed. The second phase looks how the execution time splits onto subroutines of **percolate** with different densities, to see where the bottlenecks lie concerning the execution speed.

The benchmark was performed on the Cirrus supercomputer with exclusive access to one back end node (see EPCC, 2019). The program was compiled with the GNU Fortran Compiler (**gfortran**) version 4.8.5, with the maximum optimization provided (optimization level **O3**) (see GNU Compiler Collection Team, 2019).

This paper begins by describing **percolate v0.1.0** and the conducted benchmark. Afterwards the results are presented and discussed. At last a conclusion is drawn and next steps for a follow-up benchmark are outlined.

2. Method

Let $n \in \mathbb{N}$ be a positive integer. Let $A : n \times n$ be a matrix, $A \in \mathbb{N}_0^{n \times n}$. Let $A(i, j); 1 \leq i, j \leq n; i, j \in \mathbb{N}$ be the cell of A in the i th row and the j th column. Let $S := \{filled, empty\}$ be a binary set containing the two states a cell of A can have. Let $\sigma : \mathbb{N}_0 \rightarrow S$ be a function that maps a cell of A to its state:

$$\sigma(x) = \begin{cases} filled & \text{if } x = 0 \\ empty & \text{otherwise} \end{cases}.$$

Let $A' : n + 2 \times n + 2$ be a version of A with a halo containing filled cells:

$$i, j = 0, \dots, n + 1 : A'(i, j) := \begin{cases} 0 & \text{if } i = 0 \vee j = 0 \vee i = n + 1 \vee j = n + 1 \\ A(i, j) & \text{otherwise} \end{cases}. \quad (1)$$

The density of the empty cells ρ of A is determined by the following function:

$$\rho(A) := \frac{|\{i, j = 1, \dots, n : \sigma(A(i, j)) = \text{empty}\}|}{n^2}.$$

percolate v0.1.0 randomly initializes A and sets A' according to (1) (see Algorithm 1, lines 1–2). After initialization, clusters in A' are build iteratively (see Algorithm 2). This is done by setting all empty cells of A' to their biggest neighbor. Let $\mu : \mathbb{N}_0^{m \times m} \times \mathbb{N} \times \mathbb{N} \rightarrow \mathbb{N}_0; m \in \mathbb{N}$ be the function returning the biggest value of all neighboring cells and the current cell:

$$\mu(A, i, j) := \max(A(i, j), A(i - 1, j), A(i + 1, j), A(i, j - 1), A(i, j + 1)).$$

The clustering is finished, once no cell changes value anymore (see Algorithm 2).

After the clustering the clusters are sorted in descending order based on their size ($size : \mathbb{N}_0^{n \times n} \times \mathbb{N} \rightarrow \mathbb{N}_0 : size(A, x) := |\{i, j = 1, \dots, n : A(i, j) = x\}|$)(see Algorithm 1, line 5). Sorting is done using the quicksort algorithm (see Hoare, 1961).

Afterwards every empty cell of A is mapped to its index in the array containing the sorted clusters. Every filled cell is mapped to $|clusters| + 1$. The result of this operation is the color map (see Algorithm 1, line 6). The color map is needed for writing a Portable Gray Map Image (PGM).

After doing this operations, **percolate** writes a log to **stdout** and writes A and the color map to files. These io-operations are not measured during the benchmark.

Algorithm 1 : percolate

- 1: randomly initialize A
 - 2: set A' according to (1)
 - 3: build_clusters(A')
 - 4: $A := A'(1 : n, 1 : n)$
 - 5: sort the clusters based on their size
 - 6: build the color map
-

For the benchmark, n was set to 2000. **percolate** offers an interface which controls the initialization behavior. One can define a density ρ_{goal} and a seed which is passed to the pseudo random number generator. During initialization, ρ_{goal} is approximated.

The benchmark ran **percolate** with $\rho_{goal} := 0.01, 0.02, \dots, 0.99$. Every ρ_{goal} was tried with five seeds, resulting in 495 distinct measures, with an approximately uniform distribution over ρ generated during the initialization.

Algorithm 2 : build_clusters(A')

```

1:  $A''(i, j) := 0; i, j = 0, \dots, n + 1$ 
2: while  $\sum_{i,j=1}^n (A'(i, j) - A''(i, j)) \neq 0$  do
3:    $A'' := A'$ 
4:   for  $i, j = 1, \dots, n$  do
5:      $A'(i, j) = \begin{cases} \mu(A', i, j) & \text{if } \sigma(A'(i, j)) = \text{empty} \\ 0 & \text{otherwise} \end{cases}$ 
6:   end for
7: end while

```

The three operations: clustering, sorting and building the color map were measured. The sum of all three measures is the execution time (io-operations and initialization excluded).

The execution time is measured in seconds. OpenMP 4.5's `omp_get_wtime` function was used for the time measuring (see OpenMP Architecture Review Board, 2015, Chapter 3.4.1).


3. Results

4. Discussion


5. Conclusion

References


- EPCC. Cirrus, 2019. URL <https://www.cirrus.ac.uk>.
- GNU Compiler Collection Team. The GNU Fortran Compiler, 2019. URL <https://gcc.gnu.org/onlinedocs/gcc-4.8.5/gfortran/>.
- C. A. R. Hoare. Algorithm 64: Quicksort. *Commun. ACM*, 4(7):321–, July 1961. ISSN 0001-0782. doi: 10.1145/366622.366644. URL <http://doi.acm.org/10.1145/366622.366644>.
- OpenMP Architecture Review Board. OpenMP application program interface version 4.5, 2015. URL <https://www.openmp.org/wp-content/uploads/openmp-4.5.pdf>.




(a) caption




(b) caption



(c) caption



(d) caption



(e) caption

— linear regression line

Figure 1: caption




Figure 2: Caption